

¹ dNami: a framework for solving systems of balance laws using explicit numerical schemes on structured meshes

³ **Nicolas Alferez^{*}¹, Emile Touber^{†‡^{2, 3}}, Stephen D. Winn^{2, 3}, and Yussuf Ali²**

⁵ **1** Laboratoire DynFluid, Conservatoire National des Arts et Métiers, Paris, France **2** Okinawa

⁶ Institute of Science and Technology, Okinawa, Japan **3** Department of Mechanical Engineering,

⁷ Imperial College London, London, UK

DOI: [10.21105/joss.04186](https://doi.org/10.21105/joss.04186)

Software

- [Review ↗](#)
- [Repository ↗](#)
- [Archive ↗](#)

Editor: Kristen Thyng [↗](#)

Reviewers:

- [@JamieJQuinn](#)
- [@mancellin](#)

Submitted: 10 February 2022

Published: 22 February 2022

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

⁸ Summary

⁹ The time evolution of a variety of physical and biological processes may be described by systems
¹⁰ of balance laws, which, if given appropriate initial and boundary conditions, dictate the future
¹¹ states of the systems. For instance, systems of balance laws invoking mass, momentum and
¹² energy have been incredibly successful at providing meaningful insights to the future states of
¹³ realistic systems in physics (e.g. fluid dynamics). Yet, experimenting numerically with such
¹⁴ systems still requires much implementation time. dNami (di:na:mi:) was created so that more
¹⁵ research time is spent exploring the dynamical properties of the system of balance laws of
¹⁶ interest to the user, and less time is wasted on its numerical implementation across the whole
¹⁷ computational spectrum, from the initial small-scale exploratory work on a workstation to the
¹⁸ final large-scale computations on national clusters. Thus, dNami is a computational framework
¹⁹ to study problems of the form:

$$\frac{\partial \mathbf{q}}{\partial t} = f(\mathbf{q}) + \text{initial/boundary conditions}, \quad (1)$$

²⁰ in a flexible and efficient manner, where $\mathbf{q} \in \mathbb{R}^n$ is a vector of n real-valued unknowns, t is
²¹ time, and $f(\mathbf{q})$ is a generic function of \mathbf{q} which may include differential and algebraic operators.

²² The ability of dNami to clearly separate the problem statement from its numerical imple-
²³ mentation (often a major time sink in research laboratories) is rooted in the flexibility of the
²⁴ Python language so as to let the user define her/his own system of balance laws in the most
²⁵ natural way (i.e. using a human-readable syntax), which is then interpreted in Fortran to build
²⁶ a computationally-efficient library of [Equation 1](#) which is callable from Python. Users can then
²⁷ easily interact with their own system of balance laws, including at runtime, thereby making it
²⁸ possible to integrate solutions to [Equation 1](#) with other tools and libraries (e.g. optimisation
²⁹ and stability tools) to fully explore the properties of the system, seamlessly from small to
³⁰ large-scale calculations.

³¹ State of the field

³² Code developments aimed at producing numerical solvers for [Equation 1](#) typically follow two
³³ different strategies which usually involve two different communities.

^{*}co-first author

[†]co-first author

[‡]corresponding author

34 For small-enough problems of the type of [Equation 1](#), generic solutions provided by high-
 35 level languages may be used. Notable examples may be found in the vast offer provided by
 36 the scientific Python community. One of the reasons for the Python language success in
 37 computational science is its versatility to researchers' computational needs (e.g. diverse object
 38 and data-type structures, various interactions with data at runtime using the large offer of
 39 tools from the community). However, a common major drawback is the inability to easily
 40 tackle problems involving a large number of degrees of freedom, which typically require highly
 41 efficient parallel capabilities. Note that solutions involving pre-compiled Python modules, see
 42 SciPy project ([Virtanen et al., 2020](#)), or just-in-time compilation, see Numba project ([Lam et](#)
 43 [al., 2015](#)) exist but are most of the time restricted to workstation workflow and incompatible
 44 with large-scale computing of complex problems.

45 Thus, researchers targeting [Equation 1](#) and requiring large-scale computing capabilities typi-
 46 cally have to tackle the tedious problem of High-Performance Computing (HPC) development
 47 on their own. This is time consuming for a non-specialist and often results in conservative
 48 technical solutions that do not comply with the rapid evolution of hardware architecture that
 49 comes with continuous change of parallelisation paradigms. Alternatively, such researchers can
 50 collaborate with HPC specialists for the HPC-layer of the solver. Several examples of such
 51 fruitful joint efforts are available in the literature, see for instance [Bernardini et al. \(2021\)](#)
 52 and [Krishnan et al. \(2017\)](#) in Computational Fluid Dynamics (CFD). However, this approach
 53 often results in highly technical source codes, targeting specific problems that are difficult to
 54 modify in time (e.g. new $f(\mathbf{q})$, new choice of boundary conditions).

55 A relatively recent trend aimed at maintaining HPC capabilities whilst providing user and
 56 problem-specific flexibilities is the use of Domain-Specific-Languages (DSL) libraries devel-
 57 oped by HPC specialists to tackle [Equation 1](#). Examples of such approaches can be found
 58 in CFD, a notoriously HPC-intensive domain of computational physics ([Di Renzo & Piroz-
 59 zoli, 2021](#); [Lusher et al., 2021](#); [Witherden et al., 2014](#)). Other DSL-based solvers directly
 60 target [Equation 1](#) ([Burns et al., 2020](#); [Miquel, 2021](#)). Although DSL approaches provide the
 61 versatility of the physical problem to solve and the efficient adaptability to modern hardware
 62 architectures, they do require users to learn the new DSL and drastically change paradigm in
 63 their developing approach. In addition, most DSL-based solutions rely on compiled binary ex-
 64 ecutables produced from a low-level programming language (typically C or Fortran) to achieve
 65 their performance. They do not provide the flexibility of pre-compiled Python modules from
 66 the user point of view (especially at runtime). dNami aims at conciliating the DSL-based
 67 approach with all the advantages of a Python pre-compiled module so as to solve general
 68 problems like [Equation 1](#) on both small and large scales.

69 Features

70 At the core of dNami is the translation of symbolic expressions written in high-level Python
 71 language to discretise equations in low-level Fortran language. dNami employs explicit schemes
 72 to discretise differential operators. For the temporal derivative of [Equation 1](#), a low-storage 3rd
 73 order Runge–Kutta (RK) scheme is used (other explicit schemes may easily be implemented).
 74 Spatial derivatives are discretised using finite differences of arbitrary orders provided by the
 75 user. A choice between standard and optimised schemes (such as those in [Bogey & Bailly,
 76 2004](#)) is available. Both the governing equation in the form of " $\partial\mathbf{q}/\partial t = f(\mathbf{q})$ " and the
 77 boundary conditions are specified symbolically. dNami automatically deals with stencil-size
 78 and order reduction close to boundaries using the user-specified symbolic equations, which
 79 removes the need for time-consuming and often problem-specific code development.

80 The source-to-source translation is performed by a set of Python functions using regular
 81 expressions (see `genKer.py`). The produced discretised version of [Equation 1](#) is then inserted
 82 into appropriate do-loops included in Fortran template files by pre-processing techniques. This

simple yet effective strategy makes it possible for the HPC-layer to be tailored at the template-file level independently of [Equation 1](#). Finally, the resulting Fortran source code is compiled as a shared library and optimised through the auto-optimisation process of modern Fortran compilers. A Python module is then created with a high-level interface of the shared library functions using F2PY ([Peterson, 2009](#)). It is important to stress that researchers can still follow a traditional development workflow, in a pure Fortran environment, either at the shared library level or inside the high-level interface.

[dNami](#) solves [Equation 1](#) on structured grids. Parallelisation is then ensured via classical domain decomposition techniques through point-to-point MPI communications using `mpi4py` ([Dalcin & Fang, 2021](#)). Efficient vectorisation of intense stencil-based computations are achieved via manual loop unrolling (yet automatically done by `genKer.py`) and cache-blocking techniques following recommendations from [Andreolli et al. \(2015\)](#).

Python is used at runtime to set up the run parameters and initial conditions. More importantly the time loop is exposed to Python, giving the user the freedom to interact with the computation inside the RK steps at run-time and to plug-in external libraries and/or output custom values with in-place data read and write between Python and Fortran. [dNami](#)'s Python interface thus allows easy integration of pre-processing, co-processing and post-processing tools.

101 Current [dNami](#) applications

102 [dNami](#) is currently being used by researchers to study linear and nonlinear wave phenomena
103 in various systems, from engineering to biology and geo/astro-physical flows. For example,
104 we solve the Navier–Stokes equations to study shock waves and compressible turbulence,
105 reaction-diffusion equations to study dissipative solitons in biological networks, the magneto-
106 hydrodynamic equations to study shock-entropy interactions in plasmas, and the shallow-water
107 equations to study geophysical flows (e.g. internal gravity waves, tsunamis and meteotsunamis,
108 two-dimensional compressible turbulence).

109 To illustrate the ability of [dNami](#) to quickly adapt to new systems of balance laws on supercomputers, we used the eruption of the Tonga volcano on 15 January 2022, which made
110 tsunami warnings fail in Japan ([Kataoka et al., 2022](#)). A new system of balance laws to
111 couple gravity-driven waves (shallow-water equations) in the ocean and the primary Lamb
112 wave (compressible Euler equations) in the atmosphere following the eruption was created
113 and computed on a global scale using available bathymetry data. The arrival times of the
114 meteotsunami waves are found to be in good agreement with tide gages around the globe. A
115 snapshot of the resulting animation is shown in [Figure 1](#). Starting from an empty project, the
116 global-scale meteotsunami was set and run in just a day owing to the flexibility with which
117 [dNami](#) could handle a new set of equations, a choice of spherical coordinates, and importing
118 geotiff files for the bathymetry data.

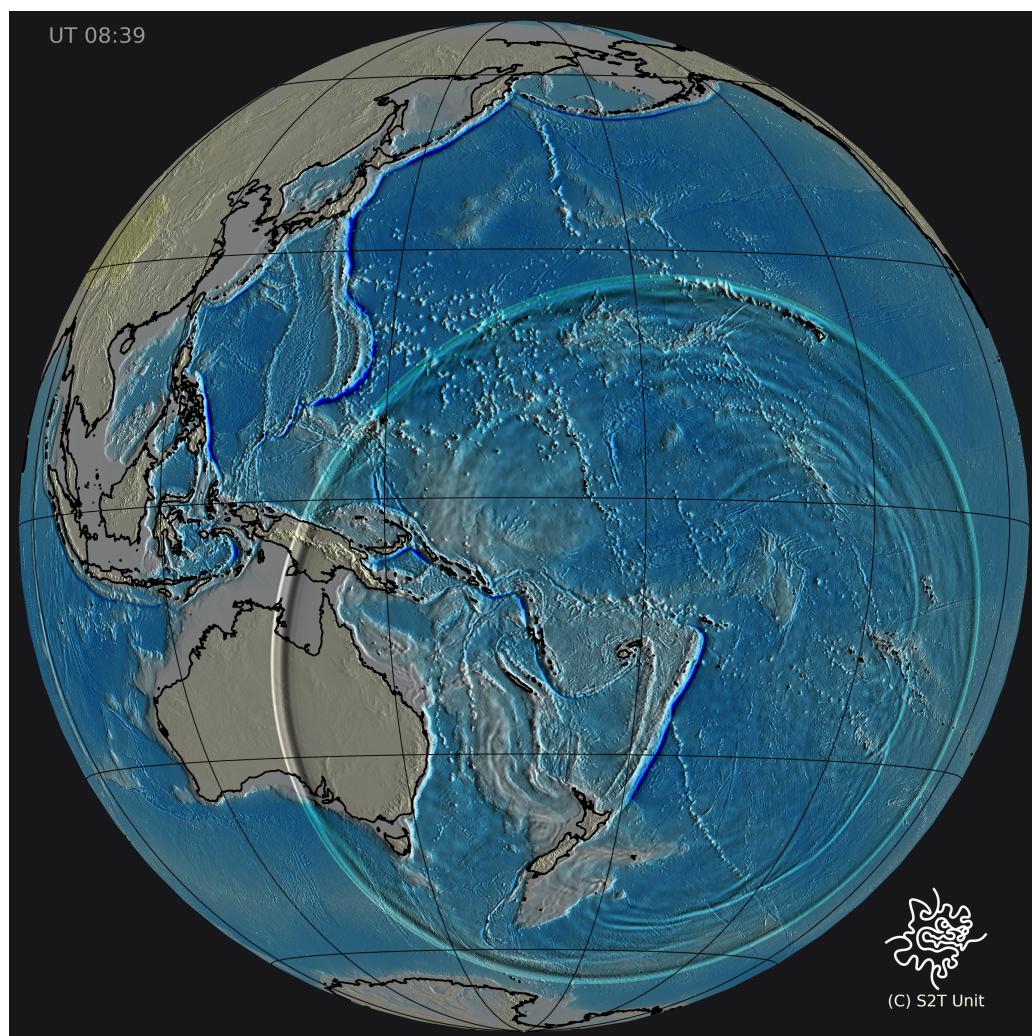


Figure 1: Example computation of a global simulation of the atmospheric and water-height disturbance due to the January 2022 Tonga volcano explosion. Time is shown in UT starting on January 15th.

120 Acknowledgements

121 We are grateful for the help and support provided by the Scientific Computing and Data Analysis
122 section of Research Support Division at OIST. This work used computational resources
123 of the supercomputer Fugaku provided by RIKEN through the HPCI System Research Project
124 (Project ID: hp200198 and hp210186).

125 References

- 126 Andreolli, C., Thierry, P., Borges, L., Skinner, G., & Yount, C. (2015). Chapter 23 - charac-
127 terization and optimization methodology applied to stencil computations. In J. Reinders &
128 J. Jeffers (Eds.), *High performance parallelism pearls* (pp. 377–396). Morgan Kaufmann.
129 <https://doi.org/10.1016/B978-0-12-802118-7.00023-6>

- 130 Bernardini, M., Modesti, D., Salvadore, F., & Pirozzoli, S. (2021). STREAmS: A high-
 131 fidelity accelerated solver for direct numerical simulation of compressible turbulent flows.
 132 *Computer Physics Communications*, 263, 107906. <https://doi.org/10.1016/j.cpc.2021.107906>
- 134 Bogey, C., & Bailly, C. (2004). A family of low dispersive and low dissipative explicit schemes
 135 for flow and noise computations. *Journal of Computational Physics*, 194(1), 194–214.
 136 <https://doi.org/10.1016/j.jcp.2003.09.003>
- 137 Burns, K. J., Vasil, G. M., Oishi, J. S., Lecoanet, D., & Brown, B. P. (2020). Dedalus:
 138 A flexible framework for numerical simulations with spectral methods. *Physical Review
 139 Research*, 2(2), 023068. <https://doi.org/10.1103/PhysRevResearch.2.023068>
- 140 Dalcin, L., & Fang, Y.-L. L. (2021). mpi4py: Status update after 12 years of development.
 141 *Computing in Science Engineering*, 23(4), 47–54. <https://doi.org/10.1109/MCSE.2021.3083216>
- 143 Di Renzo, M., & Pirozzoli, S. (2021). HTR-1.2 solver: Hypersonic Task-based Research solver
 144 version 1.2. *Computer Physics Communications*, 261, 107733. <https://doi.org/10.1016/j.cpc.2020.107733>
- 146 Kataoka, R., Winn, S. D., & Touber, E. (2022). Meteotsunamis in Japan associated with the
 147 Tonga eruption in January 2022. *EarthArXiv*. <https://doi.org/10.31223/X55K8V>
- 148 Krishnan, A., Mesnard, O., & Barba, L. A. (2017). culIBM: A GPU-based immersed boundary
 149 method code. *Journal of Open Source Software*, 2(15), 301. <https://doi.org/10.21105/joss.00301>
- 151 Lam, S. K., Pitrou, A., & Seibert, S. (2015). Numba: A llvm-based python jit compiler.
 152 *Proceedings of the Second Workshop on the LLVM Compiler Infrastructure in HPC*, 1–6.
 153 <https://doi.org/10.1145/2833157.2833162>
- 154 Lusher, D. J., Jammy, S. P., & Sandham, N. D. (2021). OpenSBLI: Automated code-
 155 generation for heterogeneous computing architectures applied to compressible fluid dy-
 156 namics on structured grids. *Computer Physics Communications*, 267, 108063. <https://doi.org/10.1016/j.cpc.2021.108063>
- 158 Miquel, B. (2021). Coral: A parallel spectral solver for fluid dynamics and partial differential
 159 equations. *Journal of Open Source Software*, 6(65), 2978. <https://doi.org/10.21105/joss.02978>
- 161 Peterson, P. (2009). F2PY: A tool for connecting fortran and python programs. *International
 162 Journal of Computational Science and Engineering*, 4(4), 296–305. <https://doi.org/10.1504/IJCSE.2009.029165>
- 164 Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D.,
 165 Burovski, E., Peterson, P., Weckesser, W., Bright, J., van der Walt, S. J., Brett, M.,
 166 Wilson, J., Millman, K. J., Mayorov, N., Nelson, A. R. J., Jones, E., Kern, R., Larson,
 167 E., ... SciPy 1.0 Contributors. (2020). SciPy 1.0: Fundamental Algorithms for Scien-
 168 tific Computing in Python. *Nature Methods*, 17, 261–272. <https://doi.org/10.1038/s41592-019-0686-2>
- 170 Witherden, F. D., Farrington, A. M., & Vincent, P. E. (2014). PyFR: An open source
 171 framework for solving advection-diffusion type problems on streaming architectures using
 172 the flux reconstruction approach. *Computer Physics Communications*, 185(11), 3028–
 173 3040. <https://doi.org/10.1016/j.cpc.2014.07.011>